



Chapter -6

Enhancing Performance with Pipelining


Pipelining is Natural: Assembly Line!

- Assuming you've got:
 - One washer (takes 30 minutes)
 - One drier (takes 40 minutes)
 - One “folder” (takes 20 minutes)



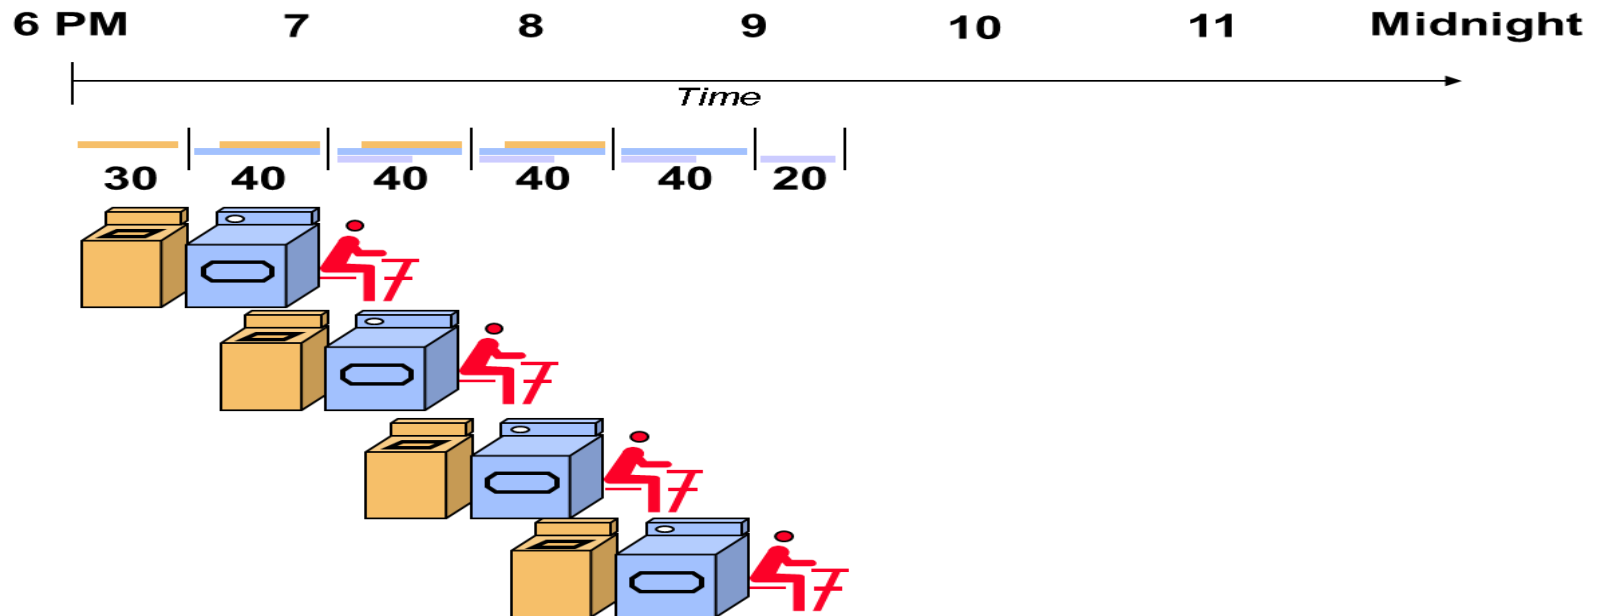
- It takes 90 minutes to wash, dry, and fold 1 load of laundry.
 - How long does 4 loads take?

The slow way(Non-pipelined)



Laundry Pipelining

- Start each load as soon as possible
 - Overlap loads



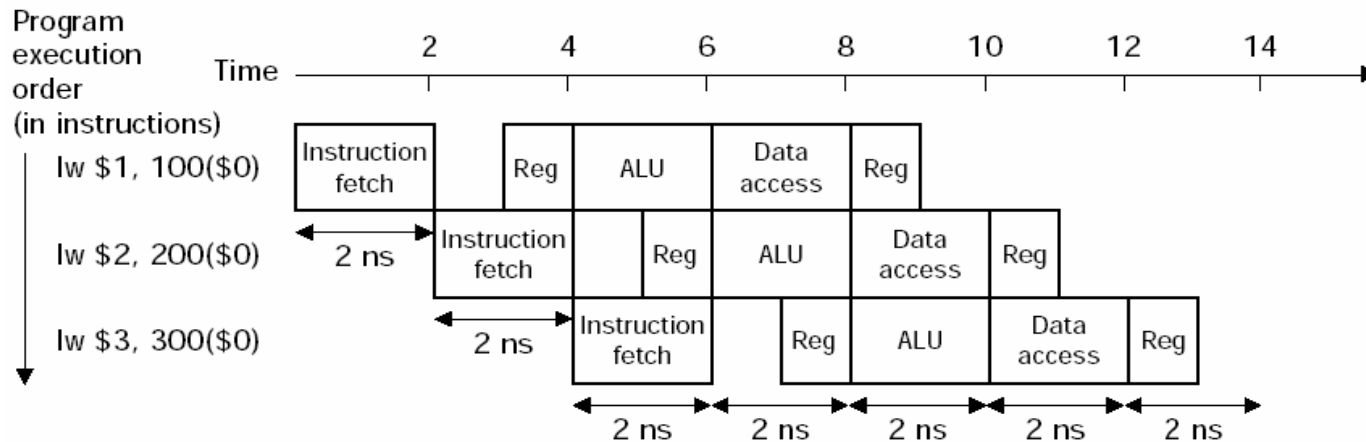
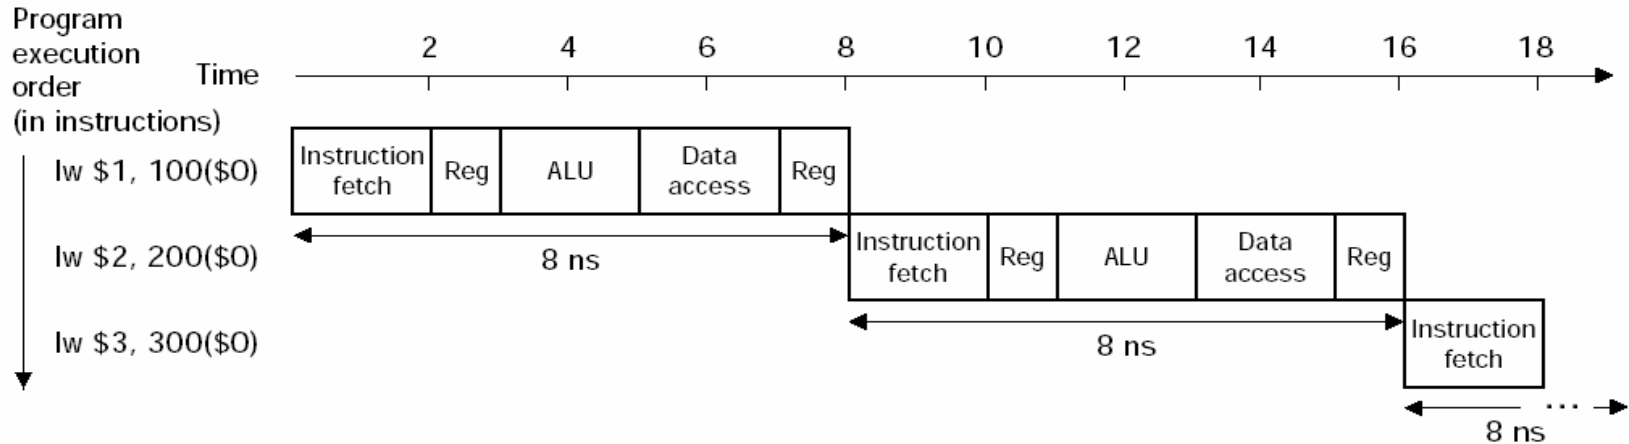
Pipelining

- Pipelining is an implementation technique in which multiple instructions are overlapped in execution.

MIPS

- Pipe Stages == The Five Execution Steps
- 1. Instruction Fetch
- 2. Instruction Decode and Register Fetch
- 3. Execution, Memory Address Computation, or Branch Completion
- 4. Memory Access or R-type instruction completion
- 5. Write-Back Step

Pipelining in MIPS



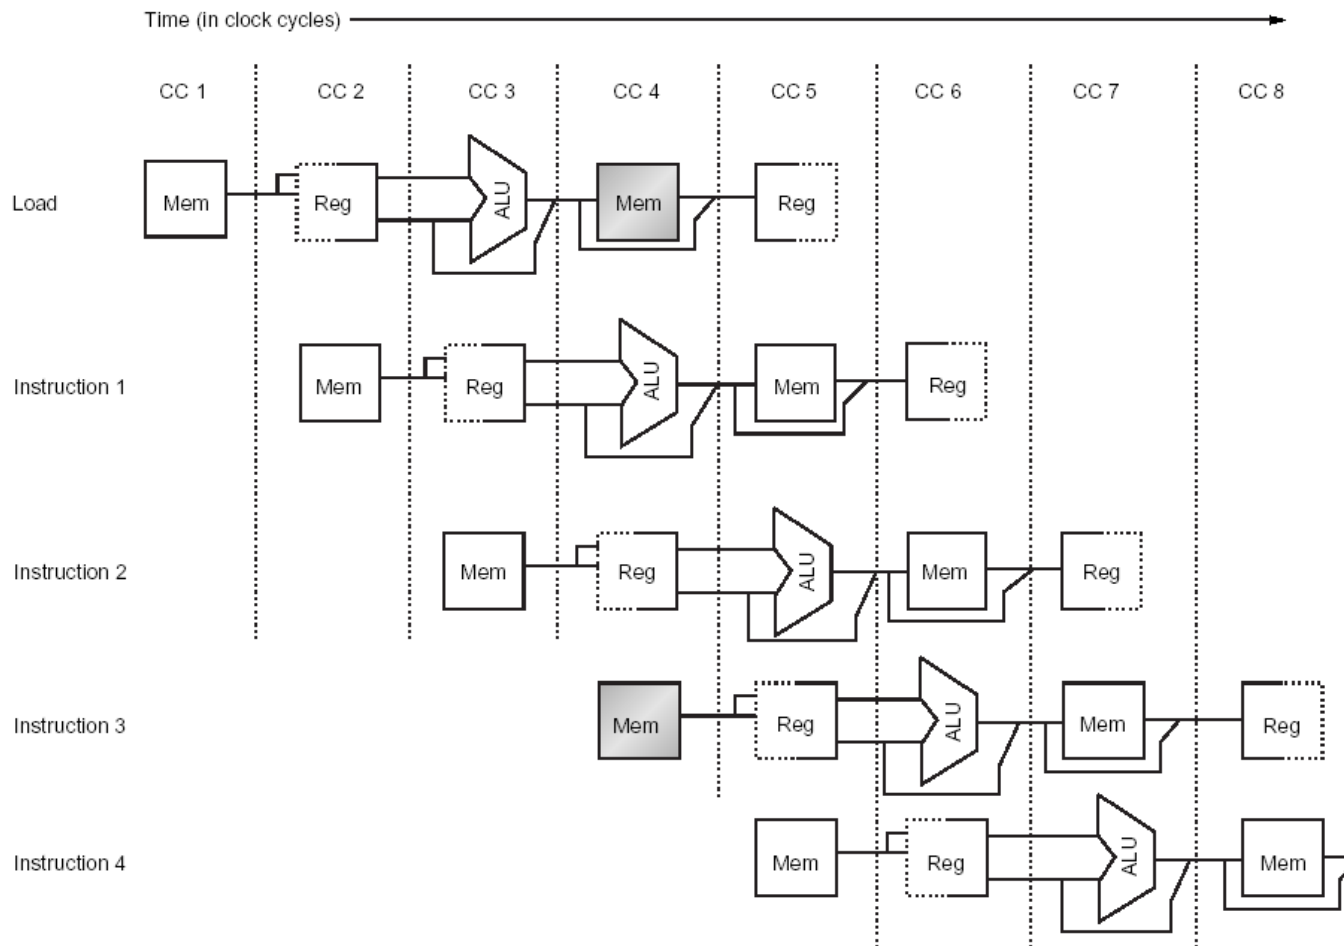
Pipeline Hazards

- There are situations in pipelining when the next instruction cannot execute in the following clock cycle. These events are called hazards.
- Three types of hazards
 - ***Structural hazards***
 - Arise from resource conflicts when the hardware can't support all possible combinations of overlapping instructions
 - ***Data hazards***
 - Arise when an instruction depends on the results of a previous instruction in a way that is exposed by overlapping of instruction in pipeline
 - ***Control hazards***
 - Arise from the pipelining of branches and other instructions that change the PC (Program Counter)

Structural Hazards

- If certain combination of instructions can't be accommodated because of resource conflicts, the machine is said to have a *structural hazard*
- It can be generated by:
 - Some functional unit is not fully pipelined
 - Some resources has not been duplicated enough to allow all the combinations in the pipeline to execute
 - For example: a machine may have only one register file write port, but under certain conditions, the pipeline might want to perform two writes in one clock cycle – this will generate structural hazard

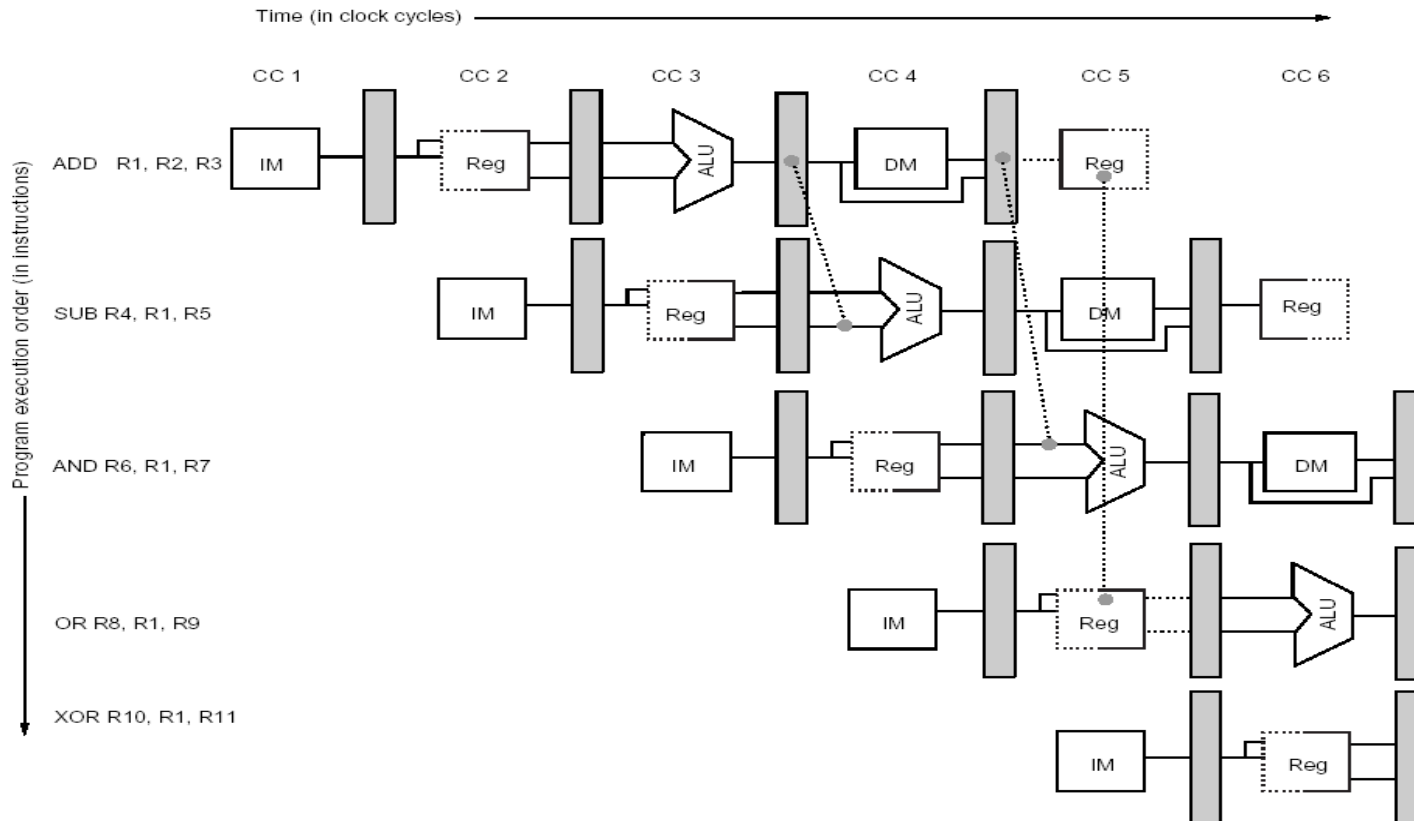
Structural Hazards(Cont)



Data Hazards

- An instruction depends on the results of a previous instruction still in the pipeline.
- Consider the execution of following instructions, on our pipelined example processor:
 - ADD R1, R2, R3
 - SUB R4, R1, R5
 - AND R6, R1, R7
 - OR R8, R1, R9
 - XOR R10, R1, R11

Data Hazards(Cont..)



Eliminate the stalls for the hazard involving SUB and AND instructions using a technique called ***forwarding***

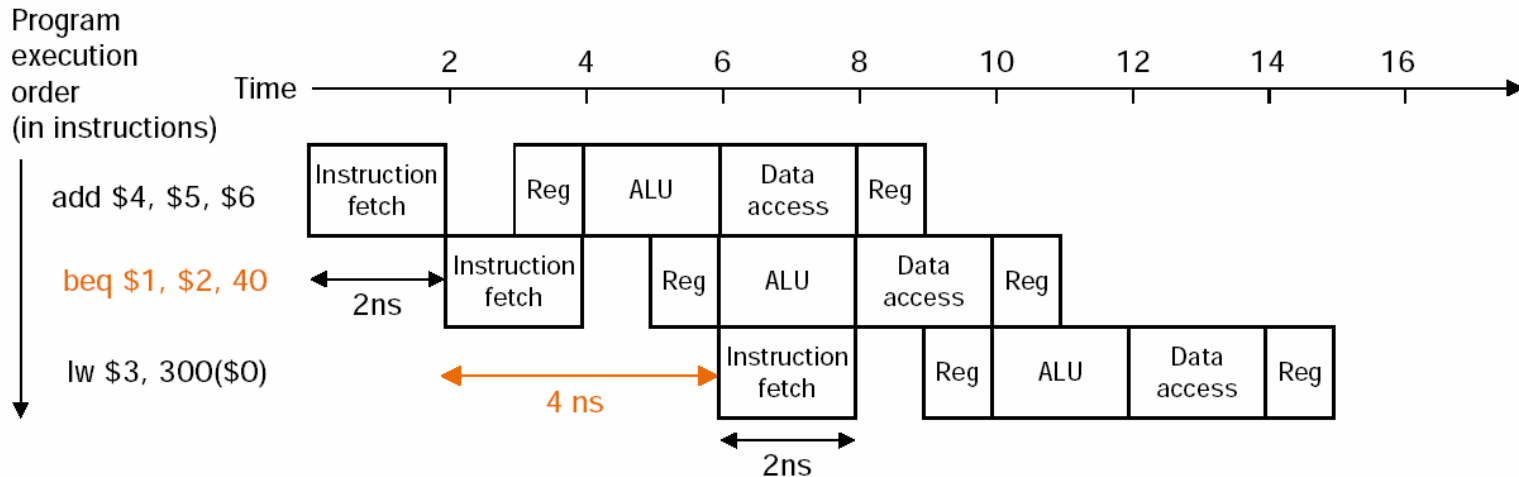
Control Hazards

- Can cause a greater performance loss than the data hazards
- When a branch is executed it may or it may not change the PC (to other value than its value + 4)
 - If a branch is changing the PC to its target address, than it is a ***taken*** branch
 - If a branch doesn't change the PC to its target address, than it is a ***not taken*** branch
- If instruction *i* is a taken branch, than the value of PC will not change until the end MEM stage of the instruction execution in the pipeline
 - A simple method to deal with branches is to stall the pipe as soon as we detect a branch until we know the result of the branch

Control Hazards

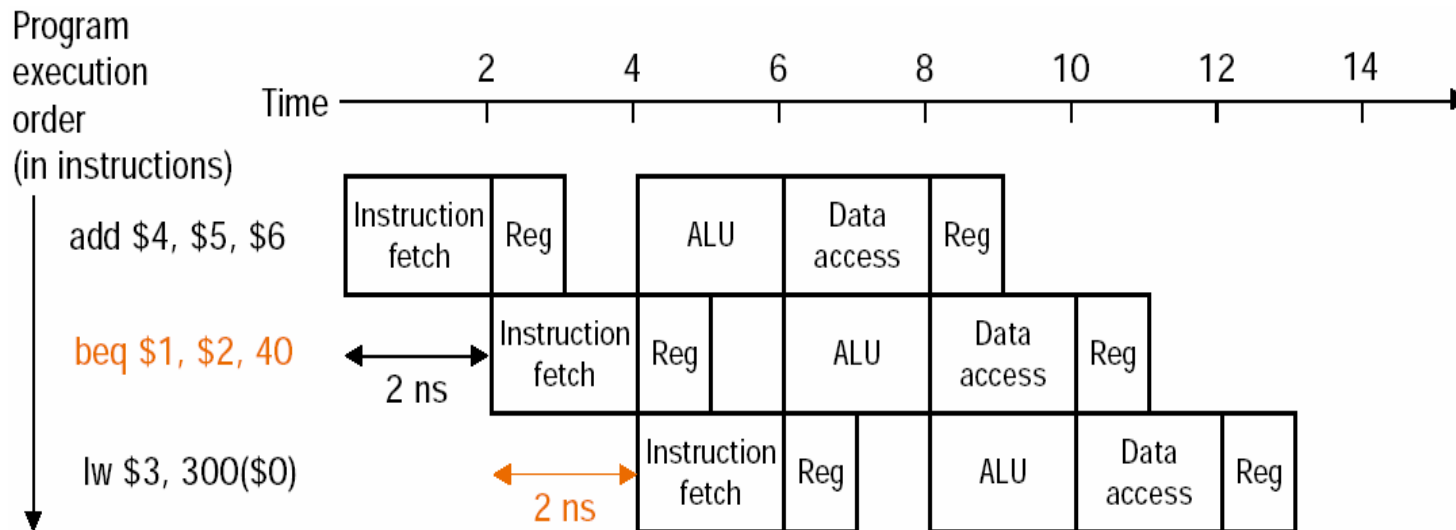
- What is the next instruction?
- Branch instructions take time to compute this.

Solution 1: Stall



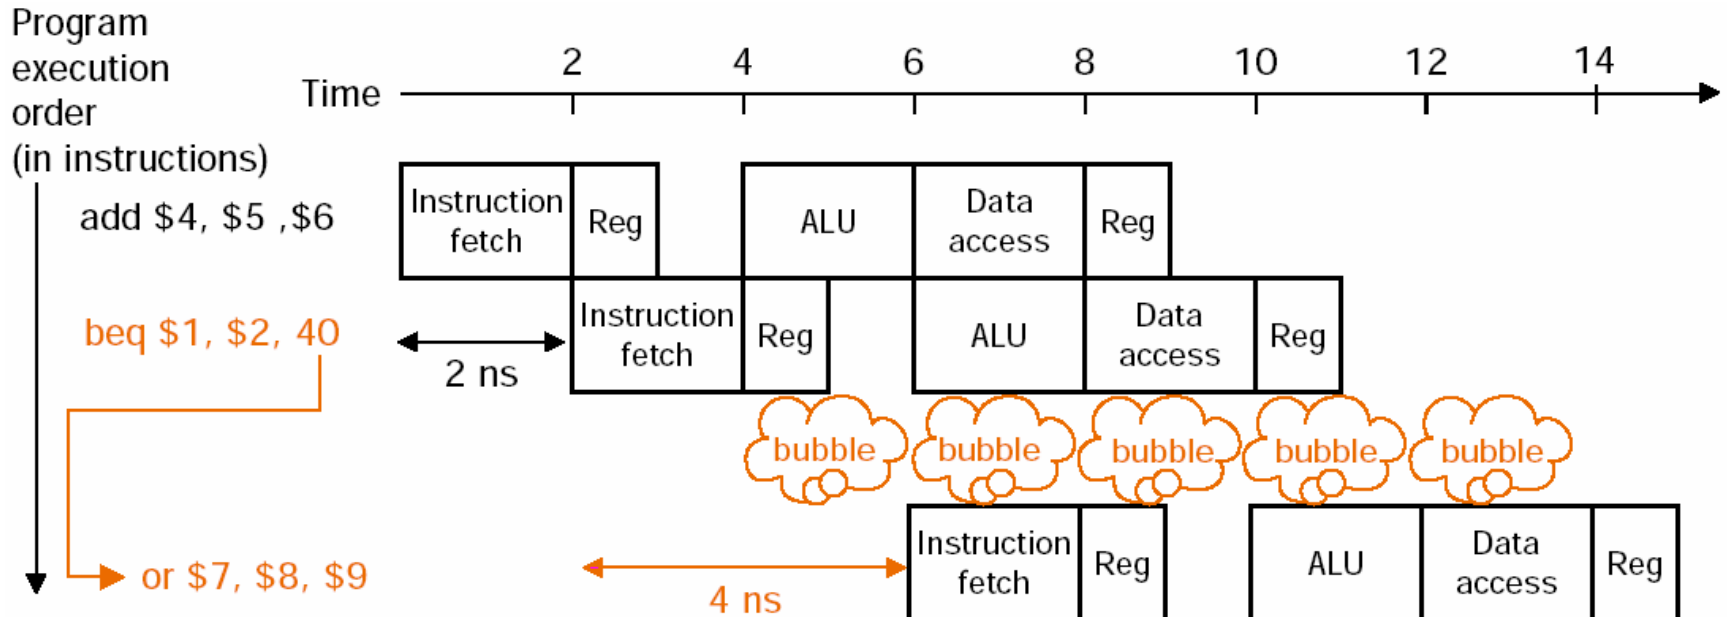
Control Hazards

- Solution 2: Predict the Branch Target



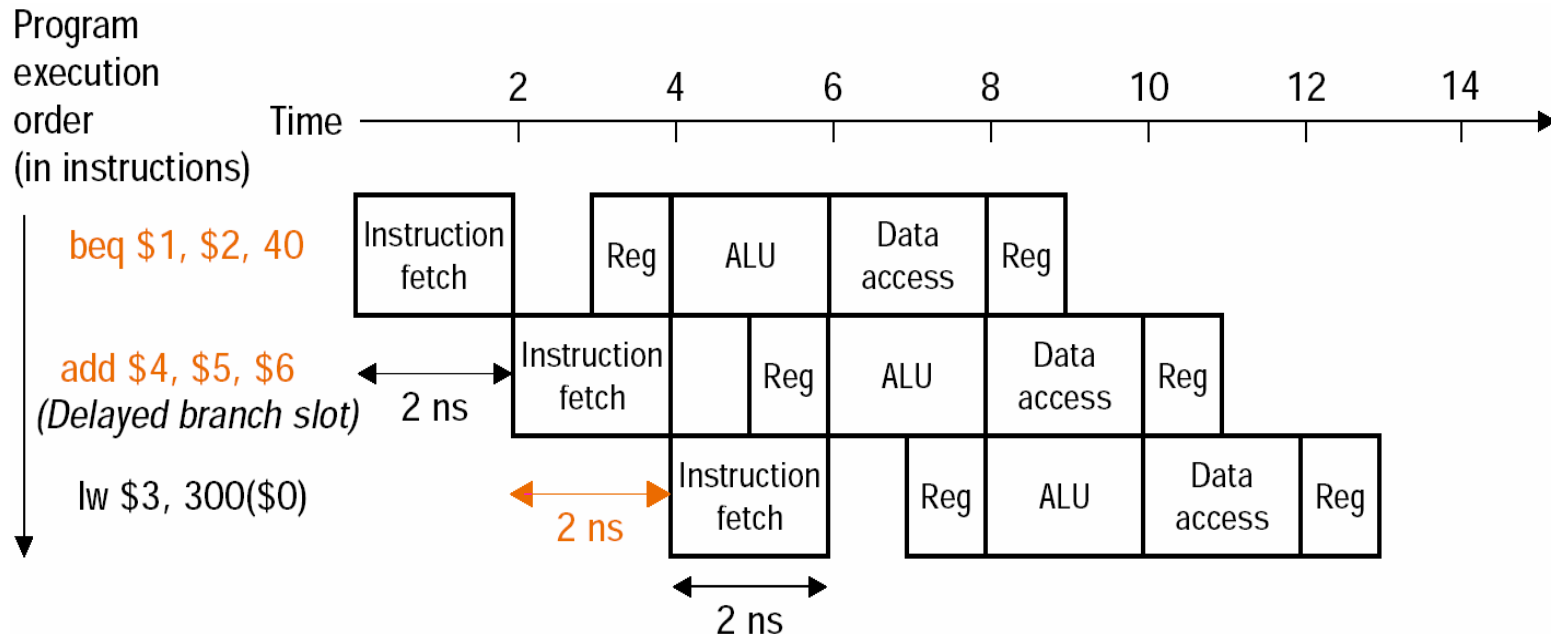
Control Hazards

Solution 2: (Mis)Predict the Branch Target



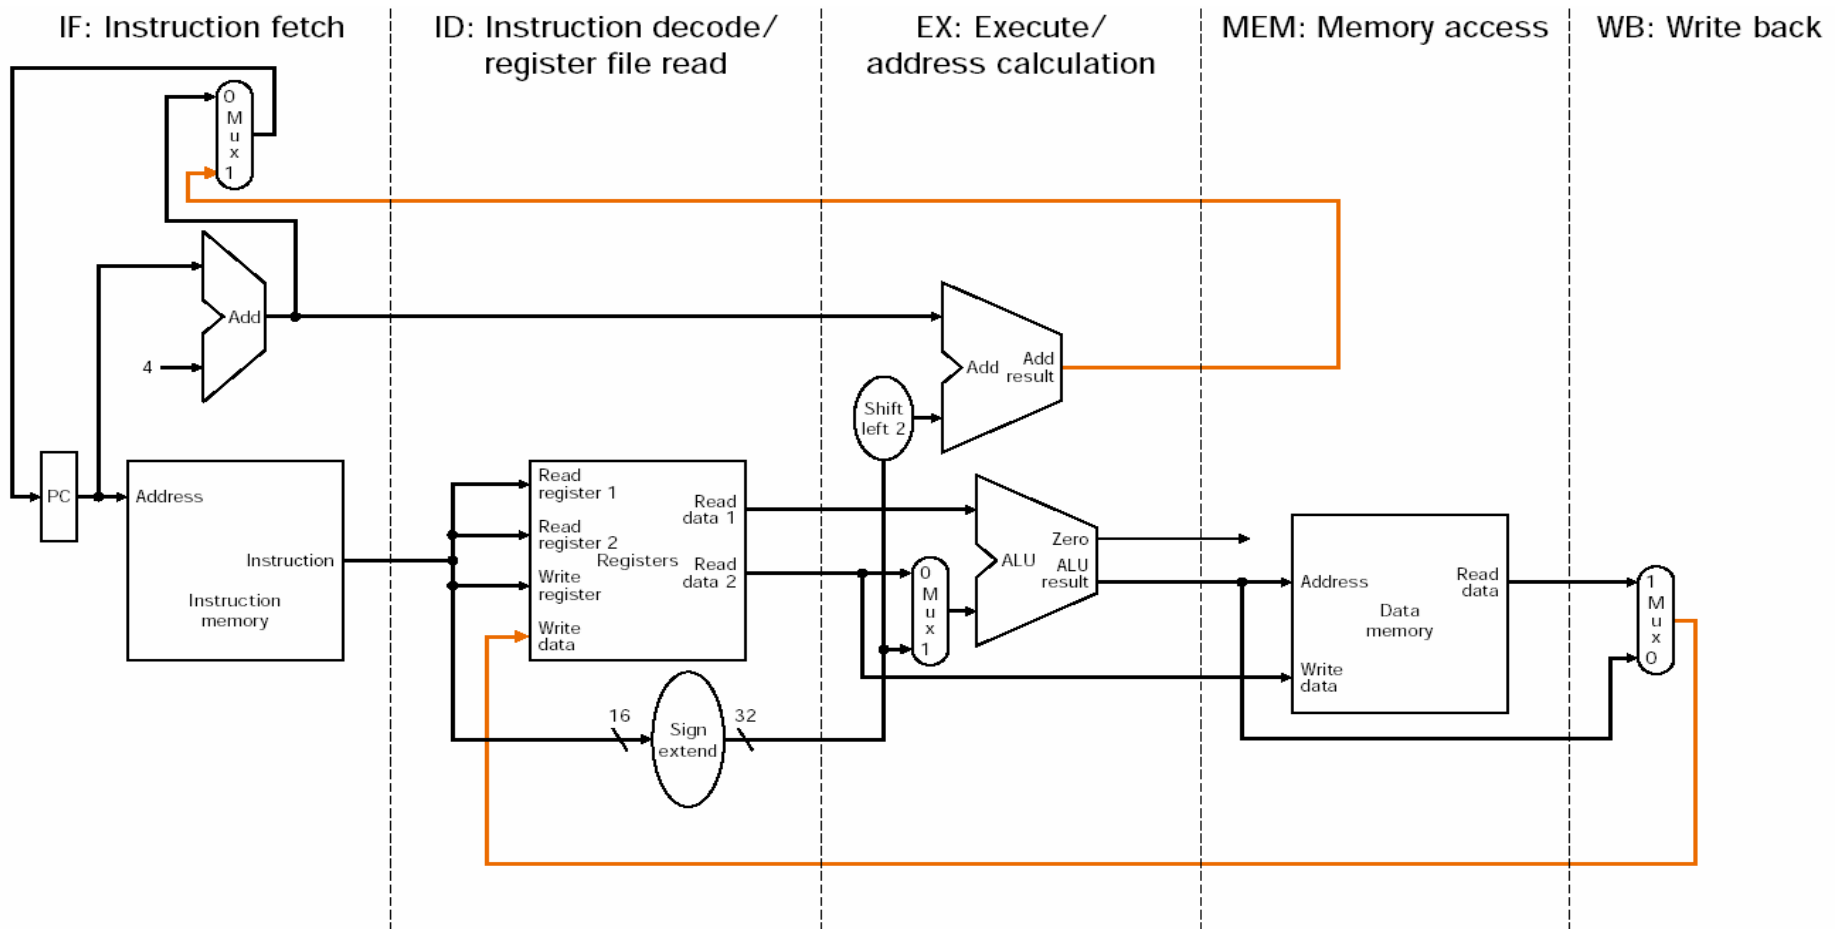
Solution 3: Delayed Decision (Used in MIPS)

- Solution 3: Delayed Decision (Used in MIPS)

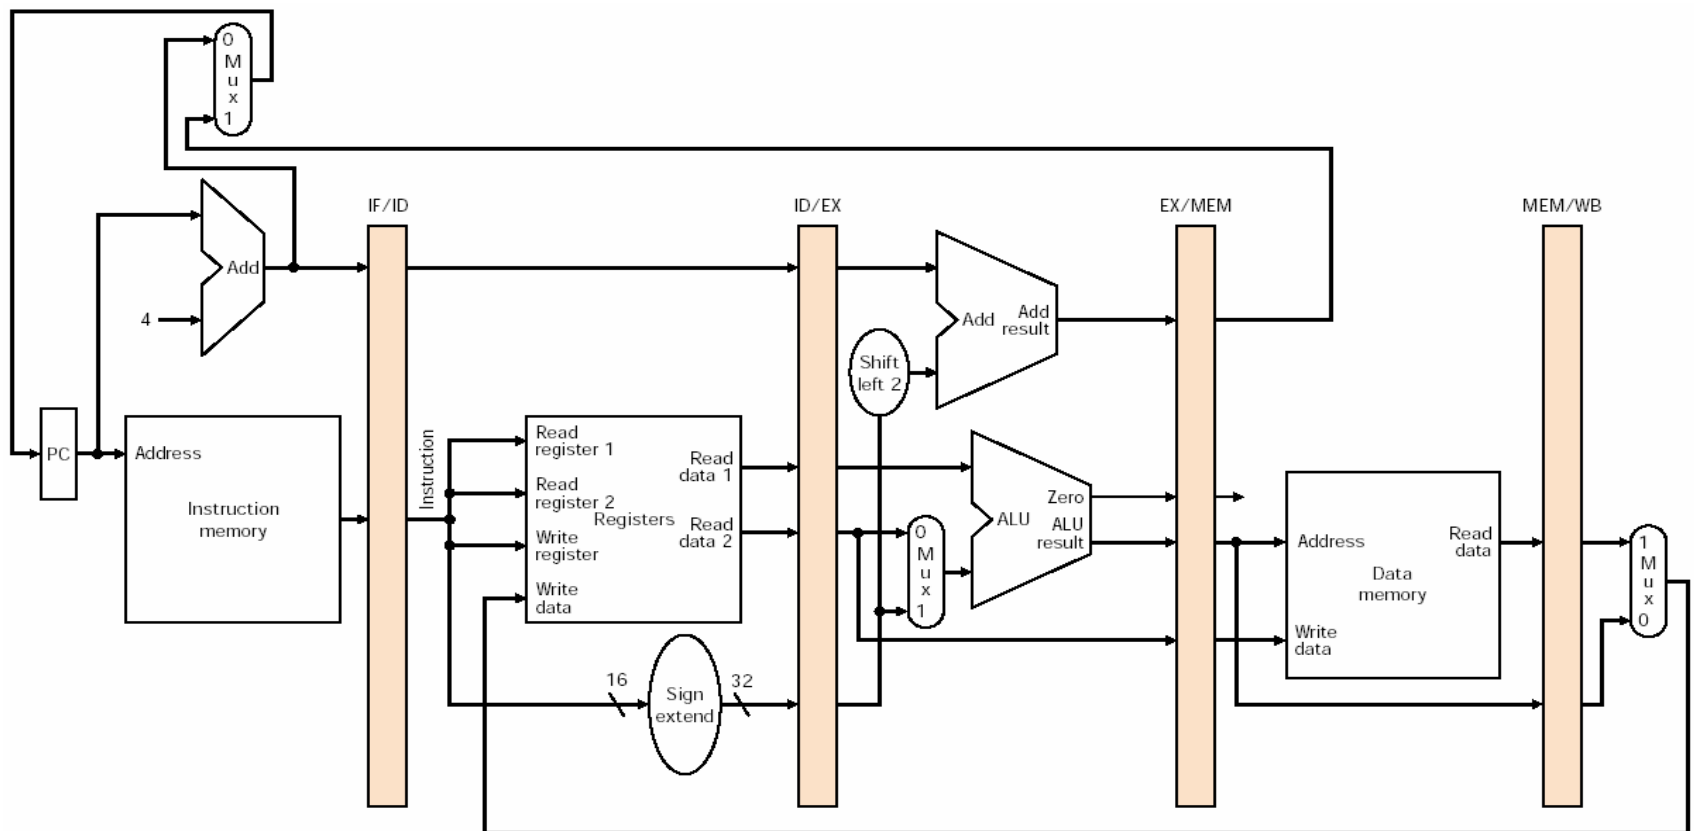


A Pipelined Datapath()

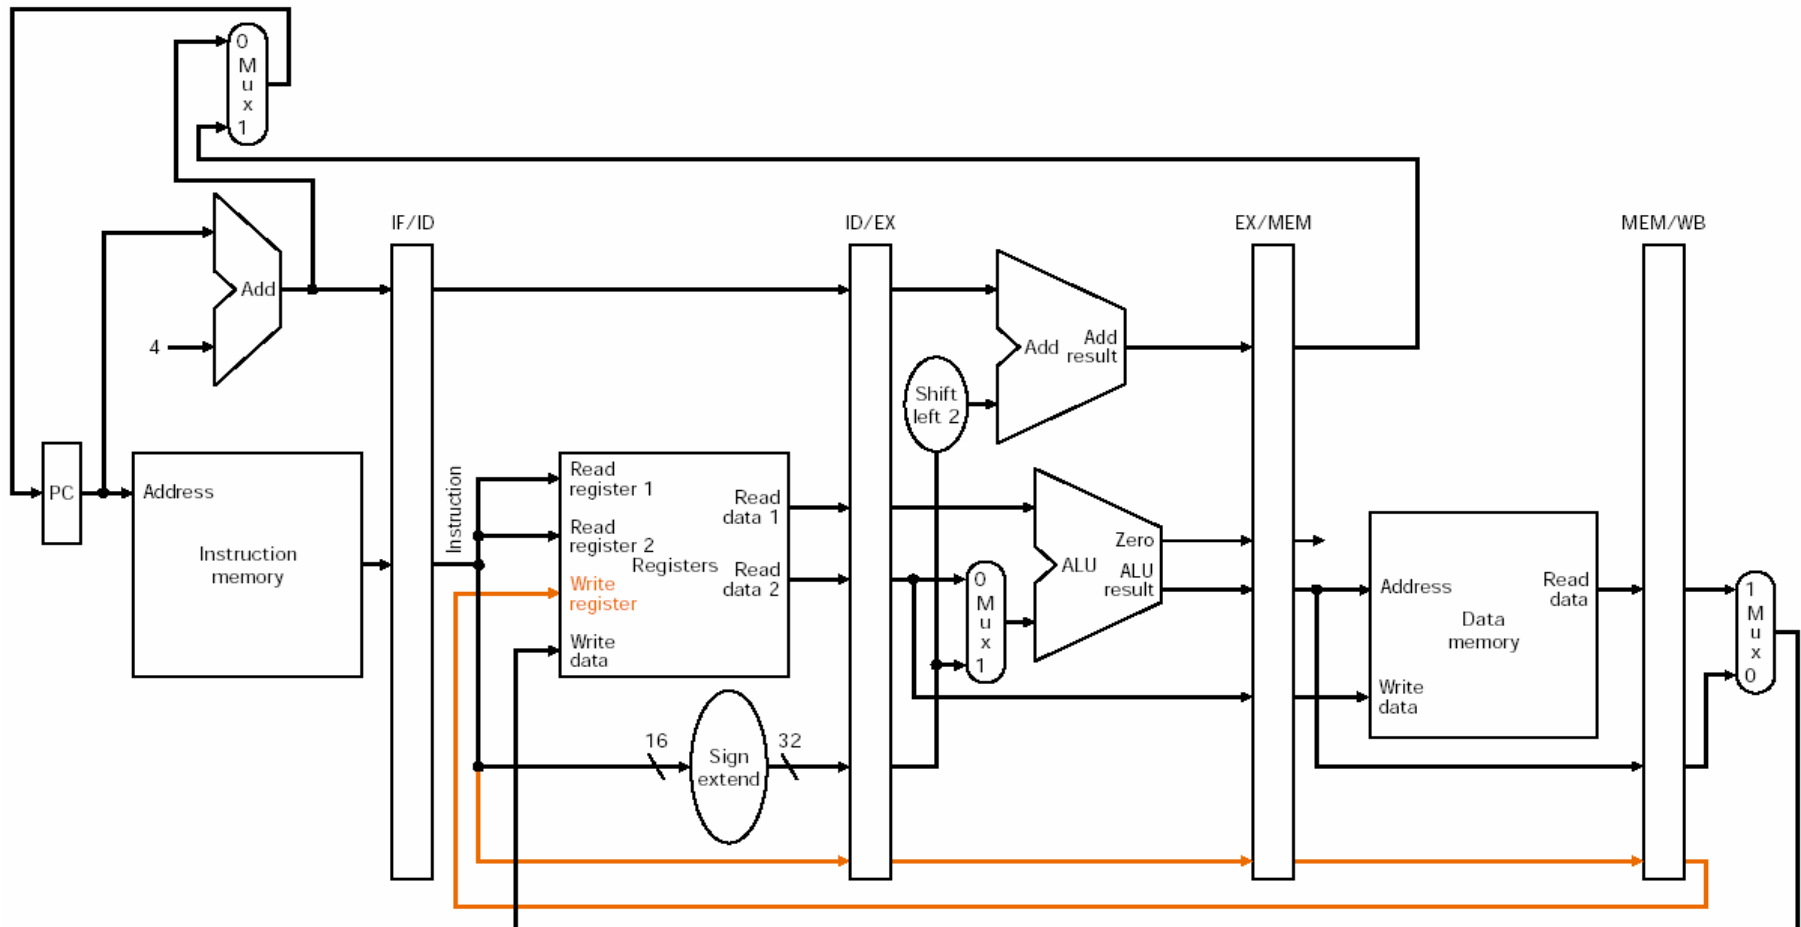
- Basic Idea



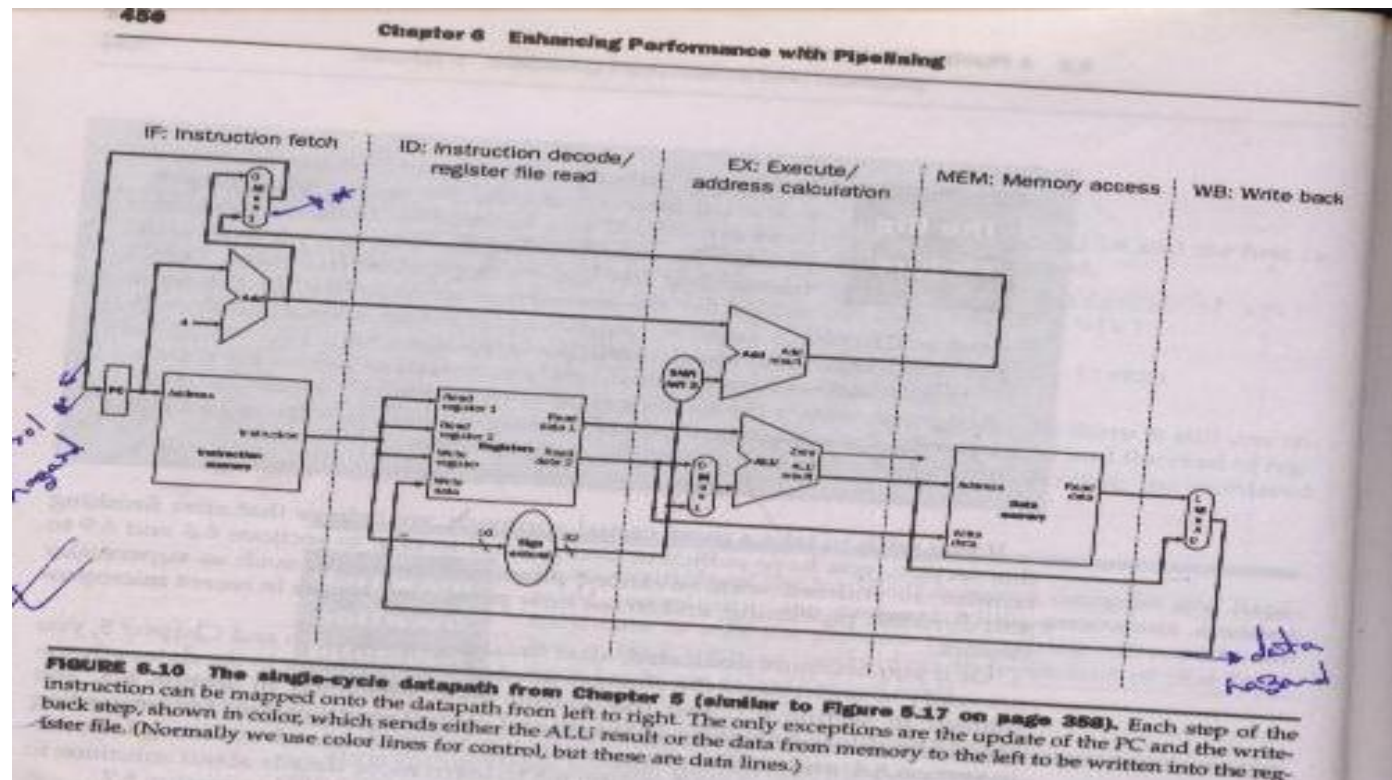
- Rectangles are pipeline registers
- Anything wrong in this picture?



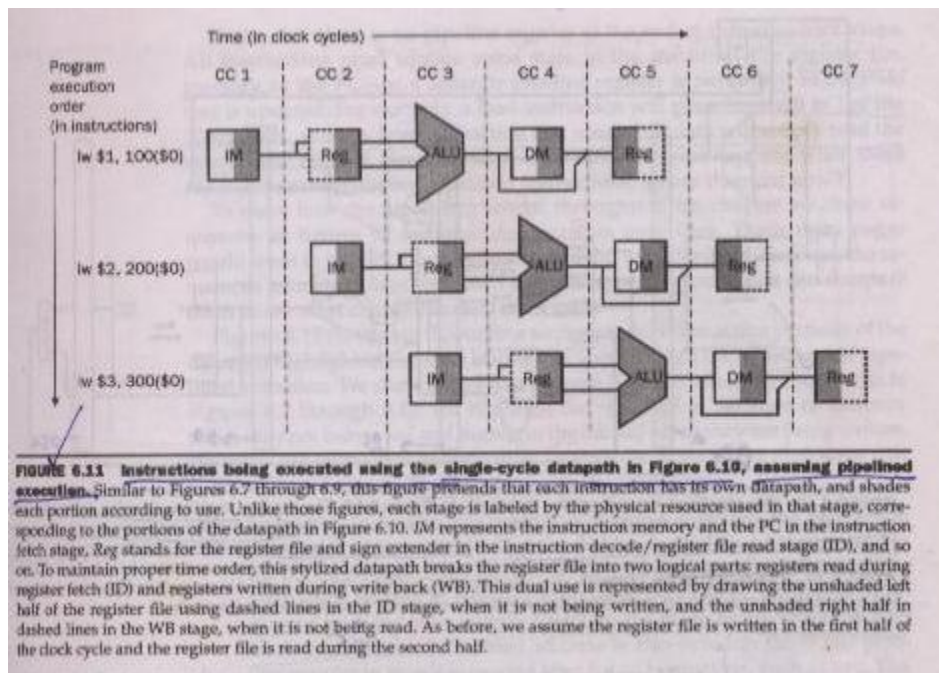
Corrected Datapath



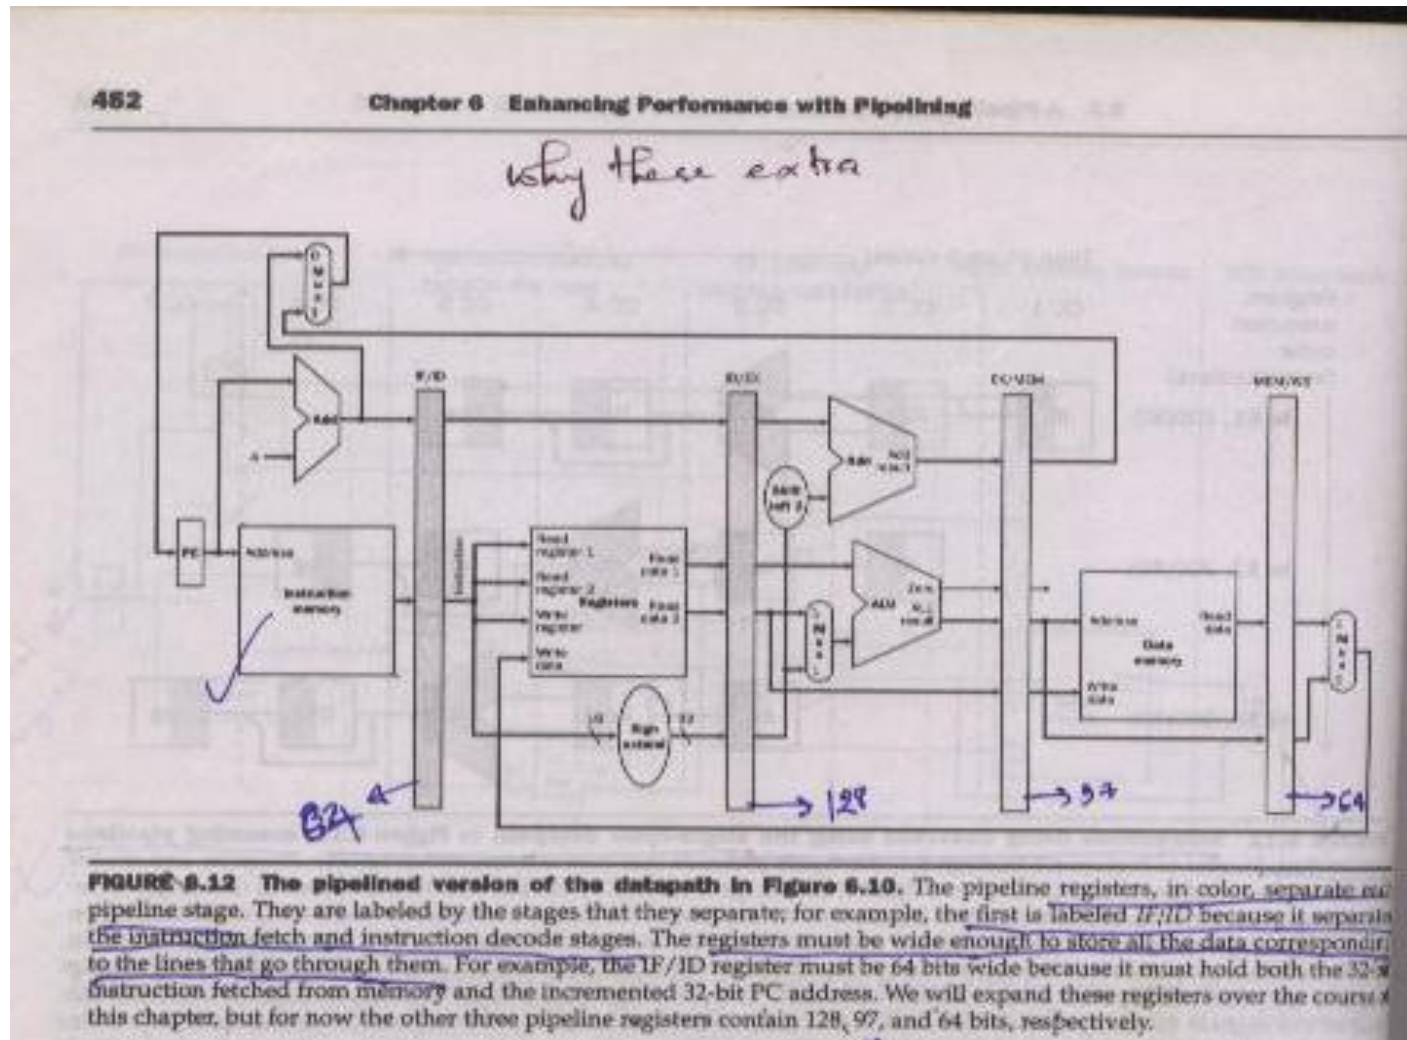
The Single Cycle Datapath



Pipeline Execution for Single Cycle Datapath



The Pipeline Version of the Datapath



The Pipeline Version of the Datapath(Cont....)

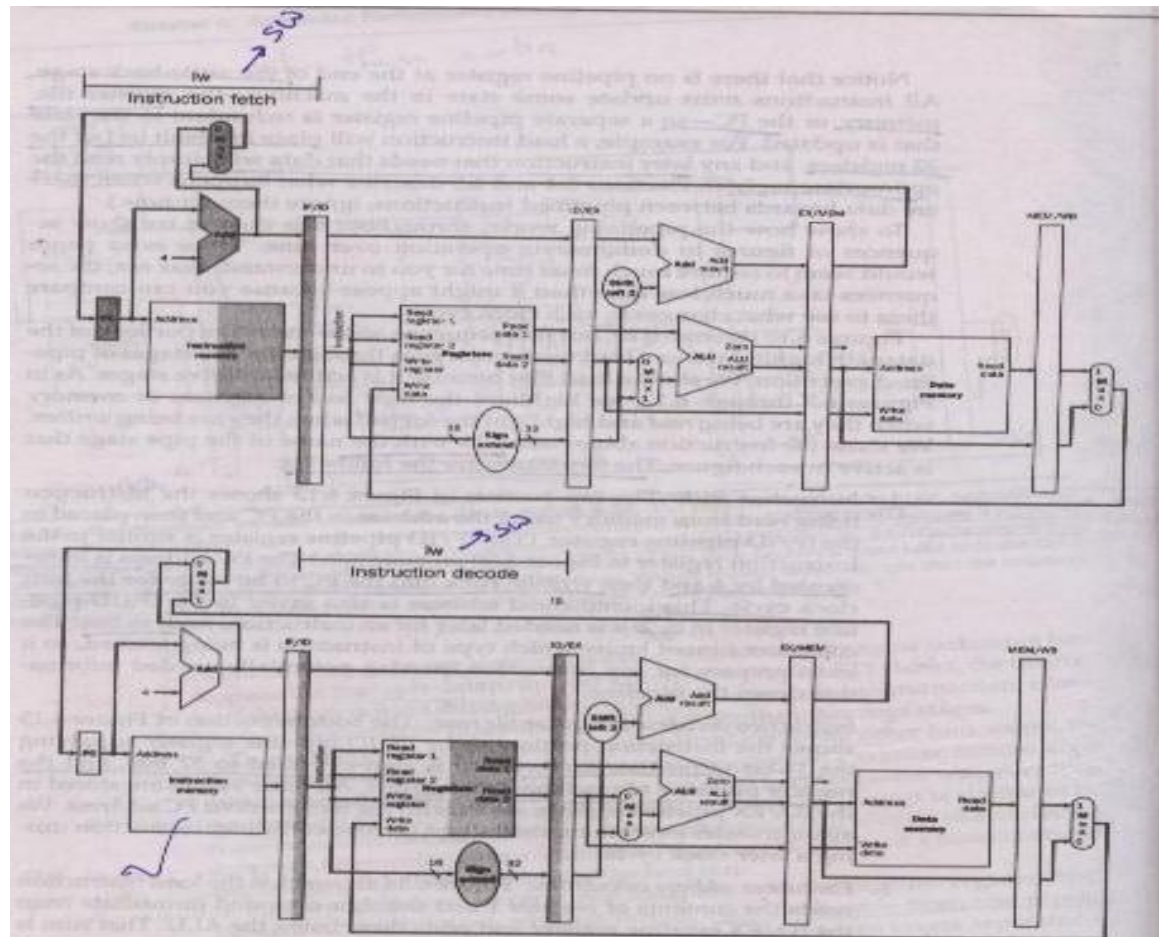
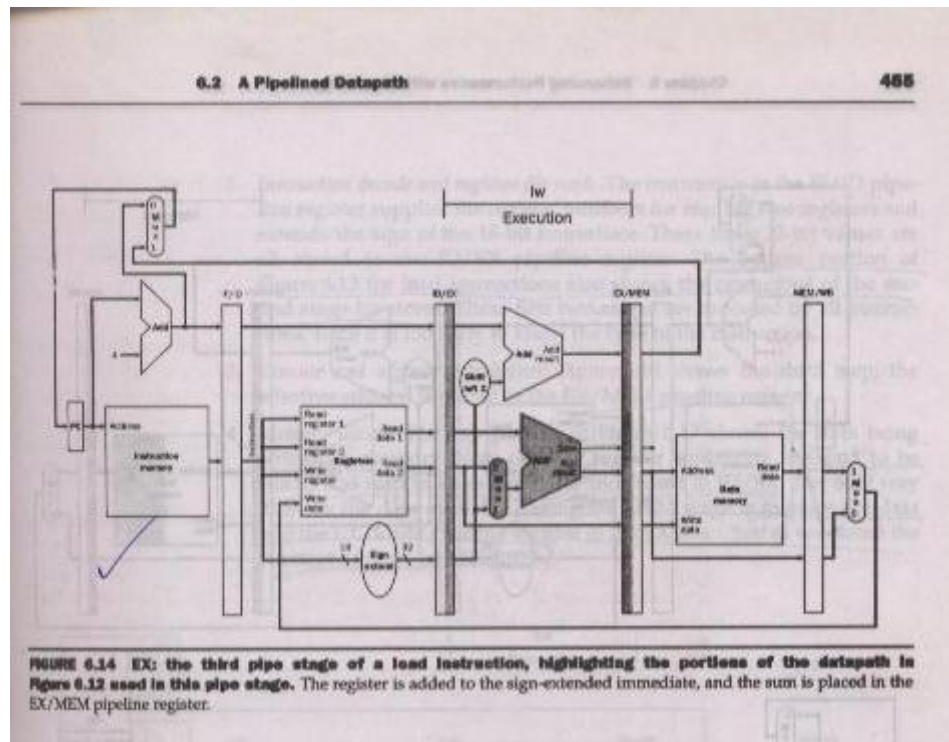
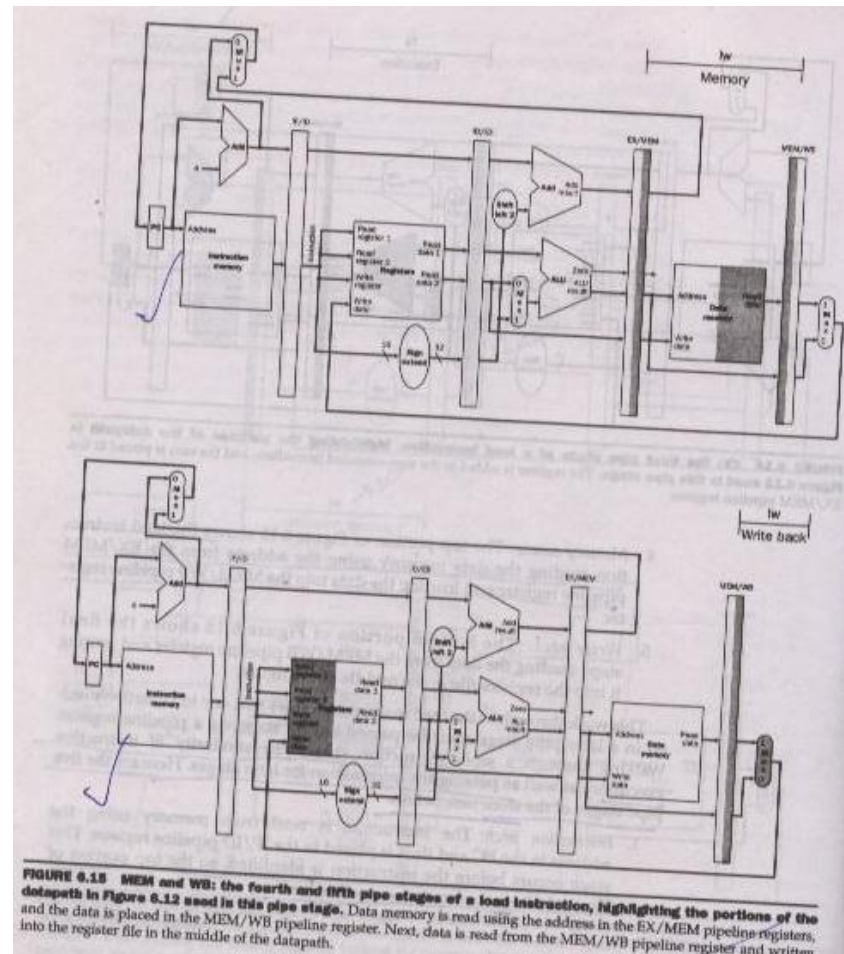


FIGURE 6.13 IF and ID: first and second pipe stages of an instruction, with the active portions of the datapath in Figure 6.12 highlighted. The highlighting convention is the same as that used in Figure 6.7. As in Chapter 5, there is no confusion when reading and writing registers because the contents change only on the clock edge. Although the load needs only the top register in stage 2, the processor doesn't know what instruction is being decoded, so it sign-extends the 16-bit constant and reads both registers into the ID/EX pipeline register. We don't need all three operands, but it simplifies control to keep all three.


The Pipeline Version of the Datapath(Cont....)



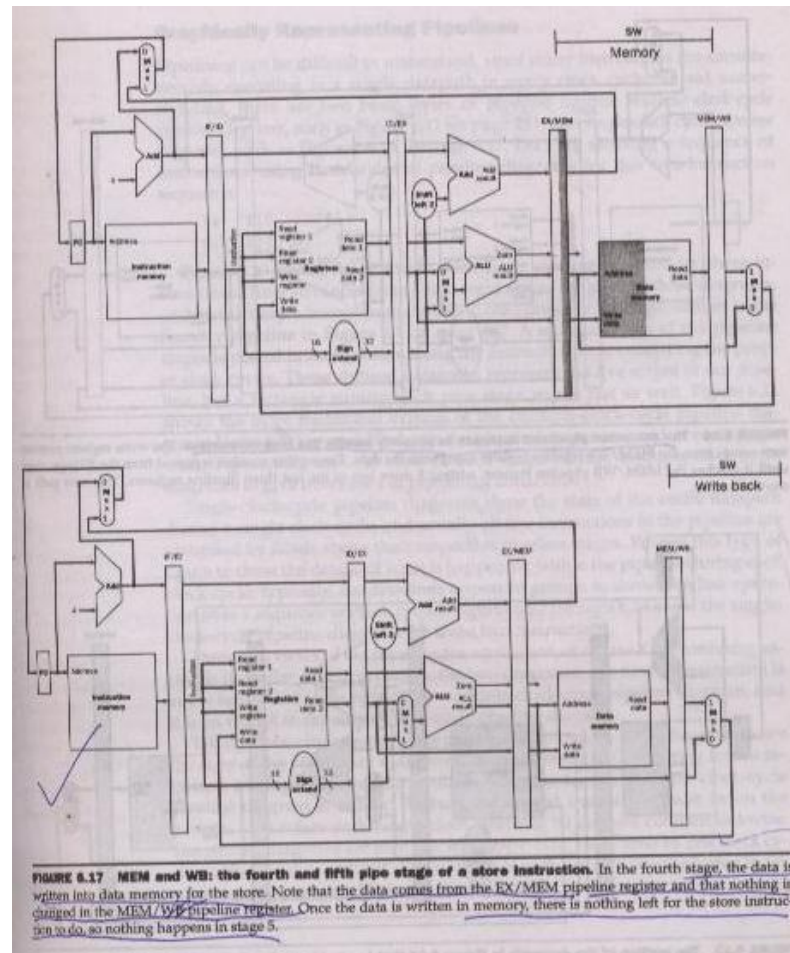
The Pipeline Version of the Datapath(Cont....)



The Pipeline Version of the Datapath(Cont....)



The Pipeline Version of the Datapath(Cont....)



Corrected pipeline version

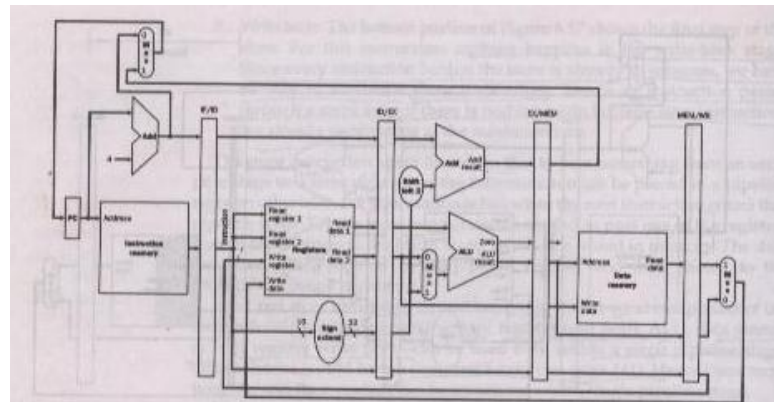


FIGURE 6.18 The corrected pipelined datapath to properly handle the load instruction. The write register number now comes from the MEM/WB pipeline register along with the data. The register number is passed from the ID pipe stage until it reaches the MEM/WB pipeline register, adding 5 more bits to the last three pipeline registers. This new path is shown in color.

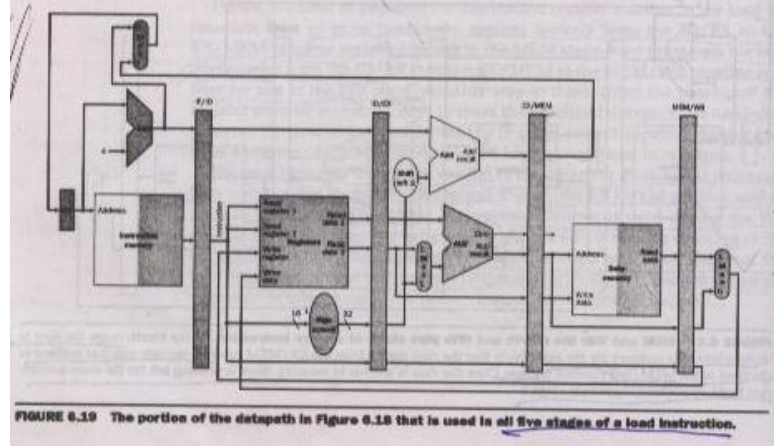



FIGURE 6.19 The portion of the datapath in Figure 6.18 that is used in all five stages of a load instruction.

Pipelined Control



Superscalar and Dynamic pipelining

- Superpipelining: Simply means longer pipeline.
- Superscaler:
- Dynamic Pipelining:

Dynamic Pipeline Scheduling

